
Dynamic Hand Gesture Recognition

Project Work

Systems Neuroscience & Neurotechnology Unit

Saarland University of Applied Sciences
Faculty of Engineering

Submitted by: Madhava Reddy Pesala

Matriculation Number: 5000804

Course of Study: Neural Engineering, MSc

Supervisor: Benedikt Buchheit, M.Eng.

Saarbrücken, September 13, 2023
Copyright c 2015 Max Mustermann, some rights reserved.

Permission is hereby granted to anyone obtaining a copy of this material, to freely copy and/or redistribute unchanged copies of this material according to the conditions of the Creative Commons Attribution-NonCommercial NoDerivatives License 4.0 International. Any form of commercial use of this material - excerpt use in particular - requires the author's prior written consent.



<http://creativecommons.org/licenses/by-nc-nd/4.0/>

Abstract

This project dives into recognizing hand movements using advanced deep learning methods. As people look for more natural ways to interact with computers and automobiles, it's becoming crucial to precisely understand hand gestures. These gestures can be used in everything from automobiles to helping those with disabilities. The heart of our project is a powerful neural network designed to pick up and understand different hand gestures from a set of videos. We enhanced our dataset with varied techniques, used a top-notch network model, and followed a detailed training plan. Our initial tests show encouraging results, matching the performance of many leading models in this field. This report takes a deep dive into the steps we took, our methods, and our findings, exciting possibilities of using deep learning for understanding hand gestures.

Declaration

I hereby declare that I have authored this work independently, that I have not used other than the declared sources and resources, and that I have explicitly marked all material that has been quoted either literally or by content from the used sources. This work has neither been submitted to any audit institution nor been published in its current form.

Saarbrücken, September 13, 2023

Madhava Reddy Pesala

Contents:

| | |
|--------------------------------------|-----------|
| 1. Introduction | 1 |
| 2. State of the Art | 1 |
| 3. Data Collection | 2 |
| 4. Preprocessing | 4 |
| 5. Methodology | 6 |
| 6. Training and Validation | 7 |
| 7. Results | 8 |
| 8. Challenges and limitations | 11 |
| 9. Future work | 11 |
| 10. References | 13 |

1 Introduction

Hand gestures have always played a major role in human-computer interaction, serving as a natural and intuitive mode of communication. With the technology we have today, the ability to recognize and interpret these gestures has opened the door to so many applications, ranging from augmented reality experiences and gaming to assistive technologies for the differently-abled. Gesture recognition, especially in our project, dynamic hand gesture recognition, involves tracking the movement of the hand over time and determining the specific gesture being made which is why it is called dynamic.

Deep learning, a subset of machine learning, has revolutionized the field of computer vision in the past decade. By neural networks with multiple layers, deep learning models can automatically learn hierarchical representations from raw data, making them particularly suitable for tasks such as image and video recognition. When applied to the domain of hand gesture recognition, deep learning techniques have shown significant promise in enhancing accuracy and robustness against various challenges such as diverse backgrounds, lighting conditions, and occlusions.

In this project, harnessing the power of deep learning to develop a dynamic hand gesture recognition system. The objective is to design a model that can effectively recognize and classify different hand gestures in real time.

2 State of the Art

Dynamic hand gesture recognition is an evolving field in computer vision and human-computer interaction. State-of-the-art in this domain have significantly improved over the past few years, with better improvements in deep learning, better datasets, and improved computing power.

Datasets:

- DHG-14/28 dataset: Contains 14 gestures and two subsets based on the number of gesture phases.
- SHREC'17 dataset: Contains 14 gestures performed with a hand's full articulation.
- Nvidia's Dynamic Hand Gesture dataset: Contains 25 gestures collected from 20 subjects in an automotive setting.

Deep Learning Models:

- 3D Convolutional Neural Networks (3D CNNs): These models process temporal sequences of spatial frames, making them suitable for dynamic gesture recognition.
- Two-Stream CNNs: Process spatial and temporal streams separately and then fuse the information.
- CNN combined with RNN/LSTM: CNNs handle the spatial features, while RNNs or LSTMs handle the temporal dynamics.
- Transformers: Although initially designed for natural language processing tasks, Transformers have been adapted for vision tasks, and their self-attention mechanisms can be useful for sequence data like dynamic gestures.

Future research is likely to focus on improving accuracy, ensuring real-time processing, handling a wider variety of gestures, and expanding the applications of dynamic hand gesture recognition.

3 Data Collection

The quality of a dataset is important for the success of any machine learning or deep learning project. In this dynamic hand gesture recognition, the dataset's quality directly influences the model's ability to accurately and consistently recognize gestures in diverse real-world scenarios. High-quality datasets contain a broad spectrum of variations, ensuring that the trained model is robust, adaptable, and less prone to errors.

For this project, I have approached two ways for dataset collection:

3.1 Generating the dataset:

In the rapidly evolving technology, state-of-the-art dynamic hand gestures often set the benchmark for user interactions. For our project, prioritized capturing these cutting-edge gestures, specifically those that have found applications in the automotive industry. The automotive sector has increasingly integrated gesture-based controls, offering drivers and passengers a seamless and advanced user experience. By focusing on these gestures, aimed to align the dataset with industry standards and address real-world applications directly. The selection of these gestures was based on their usage in modern vehicles, ensuring that our dataset remains relevant.

Detailed List of Collected Gestures:

| Gesture | Description | Contextual Information | Models |
|---|--|---|-------------------------------------|
| Swiping gesture | Swiping towards left, right, down, and up hand movements to navigate screens | Used for scrolling through options or lists | BMW, Audi, Volkswagen, Jaguar, Ford |
| Rotating gesture(Clockwise, anti-clockwise) | Circular hand movement to adjust audio volume or settings | Used for controlling volume or settings | BMW, Ford |
| Pointing gesture | Extending index finger to select an item or option | Accepting | BMW, Volkswagen, Ford, Mercedes |
| Two-finger peace gesture | Extending index and middle fingers in a peace sign gesture | Used for answering or rejecting phone calls | BMW, Mercedes Benz, Ford |
| Ok gesture | Thumb and index finger forming a circle | Used for confirming or signaling approval | Ford |

| | | | |
|------------------------------|--|--|-----|
| Thumb-pointing left gesture | Extending the thumb and pointing it toward the left side | Used for indicating forward | BMW |
| Thumb-pointing right gesture | Extending the thumb and pointing it toward the right side | Used for indicating backward | BMW |
| Pinching gesture | Bringing thumb and index finger together to zoom in or out | Used for zooming in or out on maps or screens | BMW |
| Palm wave gesture | Open hand movement to accept or dismiss notifications | Used for accepting or dismissing notifications | |
| Thumb-up gesture | Thumb raised as a sign of approval or liking | Used for approving or liking a feature | |
| Thumb-down gesture | Thumb pointed downwards to indicate disapproval | Used for rejecting or disliking a feature | |
| Pushing gesture | Forward hand movement to confirm a selection or action | Used for confirming choices or actions | |
| Pulling gesture | Backward hand movement to cancel or undo an action | Used for canceling or undoing previous actions | |

Dataset Collection for Selective Gestures Having identified the state-of-the-art dynamic hand gestures pivotal to our project, the step involves choosing selective gestures depending on our project needs. Given the expansive nature of available data and the precision required for our project, it's tough to select specific gestures prior, to ensuring that the collected data aligns seamlessly with our objectives.

To optimize the data collection process and minimize both time and effort, the next step is to approach to explore the available datasets. A comprehensive search on platforms such as Google and specialized dataset repositories will be done. The aim is to find the datasets that closely match our requirements. If we encounter datasets that match our criteria in terms of gesture type, quality, and diversity, leveraging these existing resources can provide a significant head-start. This not only minimizes the project timeline but also ensures that we're building upon tried and tested datasets. However, if the found datasets do not provide the desired results, or if there are gaps in the available data, we are prepared to supplement with our own data collection efforts to ensure comprehensiveness.

3.2 Exploring Online Repositories for Relevant Datasets

During online search for datasets, we found multiple resources. Unfortunately, access to some of these datasets was restricted, while others did not match with our specific project requirements. However, a dataset that is available on Nvidia's official website matched our requirement.

It has 25 distinct gesture classes, with data collected from a group of 20 subjects, by being conducted indoors within a car simulator. This ensures that the gestures are reflective of actual in-car interactions. SoftKinetic DS325 sensor is used to acquire frontview color and depth videos and a top-mounted DUO 3D sensor to record a pair of stereo-IR streams. This multi-faceted recording method provided left and right IR videos, color and depth videos. Each gesture captured in a video segment lasts 8 seconds. The specifics and enumeration of gestures can be found in the preprocessing section.

4 Preprocessing

Optimizing the data before model training is a good way to the success of any deep learning model. Given that the Nvidia dataset closely matches our project's requirements, we decided to use it, thereby reducing time and effort. However, even with high-quality datasets, clear checking is crucial to ensure the highest standards of results.

The Nvidia dataset offers IR, color, and depth video segments for every gesture sample. For our specific use case, since the Mediapipe framework primarily detects hands from segments with color information, extracted only the color video segments from the dataset.

It was observed that there was a similarity between some gesture classes and two gesture classes which are static to be removed as we are focussing only on dynamic gestures. To streamline our model and avoid potential bad results during recognition, merging these similar classes is performed. From the initial 25 classes, it is now reduced to 17 distinct gesture classes.

Here is the comprehensive list of gestures:

| | Initial gestures | Renamed Gestures after merging | |
|---|-------------------------|---------------------------------------|---|
| 1 | Palm swiping left | Horizontal swiping | 1 |
| 2 | Palm swiping right | | |
| 3 | Palm swiping up | Swiping Up | 2 |
| 4 | Palm Swiping up | | |
| 5 | Palm swiping down | Swiping down | 3 |
| 6 | Palm swiping down | | |
| 7 | V - Swiping left | V - Swiping left | 4 |
| 8 | V - Swiping right | V - Swiping right | 5 |

| | | | |
|----|--------------------|--------------------|----|
| 9 | V - Swiping up | V - Swiping up | 6 |
| 10 | V - Swiping down | V - Swiping down | 7 |
| 11 | Pointing | Pointing | 8 |
| 12 | Click | | |
| 13 | Pulling | Pulling | 9 |
| 14 | Pulling | | |
| 15 | Palm Opening | Palm Opening | 10 |
| 16 | Palm shake | Palm shake | 11 |
| 17 | Peace sign | Peace sign | 12 |
| 18 | Peace sign | | |
| 19 | Three fingers open | Three fingers open | 13 |

| | | | |
|----|----------------------|-----------------------------------|----|
| 20 | CW Rotation | CW Rotation | 14 |
| 21 | CCW Rotation | CCW Rotation | 15 |
| 22 | Pushing | Pushing | 16 |
| 23 | Five fingers closure | Five fingers closure | 17 |
| 24 | Thumbs up | *removed* (static gesture) | |
| 25 | OK sign | *removed* (static gesture) | |

Within each 8-second video segment of our dataset, the actual gesture performance averages around 2 seconds in duration. This indicates that a significant portion of each video, approximately 6 seconds, doesn't contribute meaningful gesture data. Such lengthy non-gesture segments can introduce noise, leading to potential misclassifications and negatively impacting the model's predictive performance. To enhance the clarity and relevance of our dataset, it's essential to trim these video segments, focusing primarily on the video frames that have hands detected. By using MediaPipe framework, the frame length of each video is trimmed to the frames that have only detected hands.

A challenge arose after trimming, the frame lengths of video segments across the entire dataset became inconsistent, complicating batch processing during model training and dimensionality issues. To address this, an average frame length across the entire dataset is identified, which amounted to 60 frames. The next step then involved standardizing all video segments to this length. Segments with frame lengths equal to or exceeding 60 were downsampled to this benchmark. Meanwhile, segments with fewer than 60 frames were excluded to maintain uniformity. Each sample was carefully checked, and any that contained inaccurate or misleading gestures were removed. Mediapipe framework is used to extract hand landmarks

from the data samples before any other preprocessing steps. This extraction resulted in clear hand landmarks set against a black background, ensuring that only the essential features were retained for training.

Following the extraction of only the color video segments from the original dataset and filtering to include only video segments with 60 frames or more, significant reduction in data samples were observed. In fact, more than half of the original data samples were excluded. This reduced dataset size was insufficient for effective model training. Given the dataset is action-based, it's crucial that any applied augmentation techniques don't disrupt the spatial and temporal information of the data. Therefore, modest rotation of about 5 degrees and translation methods were applied. These techniques were carefully chosen to enhance the dataset size without compromising the quality of the gesture representations.

With the help of Scikit-learn, the dataset is partitioned into training, testing, and validation, which constitute 70%, 15%, and 15% of the data respectively. The chosen deep learning framework for this project is PyTorch for its flexibility and efficiency. Using PyTorch's DataLoader, the datasets are loaded and batched, ensuring smooth and optimized processing during the training and evaluation phases. The batch size was set to 5, considering the available computational resources.

5 Methodology

For an action recognition dataset, which combines spatial and temporal information, the choice of the neural network model is crucial. Given the depth of deep learning models that are available, it's often good to use pretrained models, not only to lessen the training process but also to make use of the knowledge from vast datasets. The Convolutional Neural Network (CNN) good at extracting spatial features, while the Recurrent Neural Network (RNN) at capturing temporal sequences. With this understanding, several models rise for consideration: ResNet3D, CNN+RNN, 3D CNN, C3D, and the relatively new Transformers, which have shown promise in sequence-to-sequence tasks.

In the field of deep learning, there's a level of unpredictability. Often, the best model for a specific task isn't immediately evident until used into trial and experimentation. So, it is decided to go with the ResNet3D pretrained model at first. This model has been trained on the Kinetics dataset, a large-scale, high-quality dataset of YouTube video URLs, which covers a wide range of human actions. The Kinetics dataset encompasses around 400 action classes with at least 400 video clips for each action, each 10 seconds long, making it a robust choice for initial transfer learning models. The richness of this dataset, combined with the depth of the ResNet3D architecture, offers a solid foundation for recognizing dynamic hand gestures. The ResNet architecture, with its unique residual connections, effectively for the vanishing gradient problem. To be precise, ResNet3D-18 architecture was chosen for this project, which gives a balance between computational efficiency and the size of the data we have.

6 Training and Validation

Device Configuration:

The computational device is selected based on the availability of CUDA. If CUDA is available, it makes that a GPU is present and can be utilized for computation, ensuring faster training and validation. If not, computations default to the CPU.

Loss Function:

Cross Entropy Loss is selected as the loss criterion for our model. This decision is particularly apt for classification tasks, given that this loss function quantifies the difference between the predicted probabilities output by the model and the actual class labels of the data. In the context of gesture recognition, the use of Cross Entropy Loss ensures that the model is penalized appropriately for incorrect predictions, making it to produce more accurate results over subsequent training iterations.

Optimizer:

Stochastic Gradient Descent (SGD) is Chosen for its robustness and widespread use in training deep neural networks. Specific parameters include: Learning Rate of 0.0005, dictates the step size during optimization. Momentum of 0.9, helps accelerate the optimizer towards the correct direction and dampens oscillations and this value is commonly used. Weight Decay of 0.00001 Serves as a regularization technique, preventing the model from fitting too closely to the training data by penalizing large coefficients.

Learning Rate Scheduler:

The ReduceLROnPlateau scheduler is deployed, aimed at decreasing the learning rate when the model's performance plateaus. If the validation loss doesn't show improvement for a set number of epochs (4 epochs in this case), the learning rate is reduced by a factor of 0.1. This strategy aids in achieving a more refined model convergence.

Training Loop:

The model undergoes training for 35 epochs. In each epoch, the model is set to training mode. Data is fetched batch-wise from the training data loader. The forward pass computes the model's predictions, followed by the calculation of the loss using the pre-defined criterion. Backpropagation computes the necessary gradients, and the optimizer updates the model parameters accordingly. The training accuracy is computed for the epoch by comparing the model's predictions against the true labels. Post-training in each epoch, the model's performance is validated using the validation dataset. The model transitions to evaluation mode, ensuring that certain layers like dropout behave differently than during training. Validation loss and accuracy metrics are determined, providing insights into the model's performance on unseen data.

Fine Tuning:

In the preliminary stages of the project, typical default values were employed for hyperparameters such as batch size, learning rate, and weight decay. However, as the training process evolved, several rounds of hyperparameter tuning were performed to optimize model performance.

One significant constraint encountered was the limitation of computational resources. This restricted the batch size to a maximum of 5, a value that typically might be increased to capitalize on parallel processing capabilities and achieve smoother gradient descent.

As the training progressed, it indicated overfitting, that the model was too closely fitting the training data by significant difference with validation data, which might perform poorly on unseen data. To mitigate this, regularization techniques were explored. L2 regularization, a form of weight decay, was incorporated to constrain the network's weights, discouraging overly complex models that might overfit. Furthermore, early stopping was introduced, providing a mechanism to halt training if the model's validation performance ceased to improve, preventing potential overfitting.

Upon inspecting the model's architecture, the absence of dropout layers was noticed. Dropout, which randomly deactivates a subset of neurons during training, was then introduced as an additional measure against overfitting. However, this appeared causing slight underfitting. Recognizing this, the dropout layers were subsequently removed to strike a balance.

Due to the computational limitations, significant alterations in hyperparameters were challenging. The batch size, for instance, remained at 5 throughout the experimentation phase.

7 Results

Regarding the results, due to limited computational resources, for the time being, the model was trained on just 5 out of the 17 gesture classes, and only the rotation augmentation technique was employed. However, the final code provided has both rotation and translation augmentation techniques and is designed to accommodate all 17 gesture classes.

Selected gestures were: {Palm Opening, Palm Shake, Pulling, Pointing, CCW_Rotation}

```

Epoch [1/35] - Train Loss: 1.5130, Train Accuracy: 34.44%, Validation Loss: 1.3295, Validation Accuracy: 44.33%
Epoch [2/35] - Train Loss: 1.3001, Train Accuracy: 48.89%, Validation Loss: 1.0082, Validation Accuracy: 68.04%
Epoch [3/35] - Train Loss: 1.0628, Train Accuracy: 60.89%, Validation Loss: 0.7057, Validation Accuracy: 74.23%
Epoch [4/35] - Train Loss: 0.7775, Train Accuracy: 76.00%, Validation Loss: 0.4878, Validation Accuracy: 90.72%
Epoch [5/35] - Train Loss: 0.5857, Train Accuracy: 81.78%, Validation Loss: 0.4724, Validation Accuracy: 84.54%
Epoch [6/35] - Train Loss: 0.4691, Train Accuracy: 89.33%, Validation Loss: 0.4448, Validation Accuracy: 89.69%
Epoch [7/35] - Train Loss: 0.3266, Train Accuracy: 92.22%, Validation Loss: 0.2868, Validation Accuracy: 89.69%
Epoch [8/35] - Train Loss: 0.3518, Train Accuracy: 89.78%, Validation Loss: 0.3126, Validation Accuracy: 89.69%
Epoch [9/35] - Train Loss: 0.2358, Train Accuracy: 94.00%, Validation Loss: 0.2623, Validation Accuracy: 90.72%
Epoch [10/35] - Train Loss: 0.2217, Train Accuracy: 95.11%, Validation Loss: 0.3411, Validation Accuracy: 91.75%
Epoch [11/35] - Train Loss: 0.1323, Train Accuracy: 98.22%, Validation Loss: 0.1452, Validation Accuracy: 95.88%
Epoch [12/35] - Train Loss: 0.1471, Train Accuracy: 96.22%, Validation Loss: 0.1383, Validation Accuracy: 95.88%
Epoch [13/35] - Train Loss: 0.0991, Train Accuracy: 98.22%, Validation Loss: 0.0949, Validation Accuracy: 96.91%
Epoch [14/35] - Train Loss: 0.1289, Train Accuracy: 97.11%, Validation Loss: 0.1488, Validation Accuracy: 93.81%
Epoch [15/35] - Train Loss: 0.1123, Train Accuracy: 98.00%, Validation Loss: 0.2819, Validation Accuracy: 90.72%
Epoch [16/35] - Train Loss: 0.0719, Train Accuracy: 99.33%, Validation Loss: 0.0903, Validation Accuracy: 97.94%
Epoch [17/35] - Train Loss: 0.0534, Train Accuracy: 99.11%, Validation Loss: 0.0640, Validation Accuracy: 100.00%
Epoch [18/35] - Train Loss: 0.0656, Train Accuracy: 98.89%, Validation Loss: 0.0890, Validation Accuracy: 97.94%
Epoch [19/35] - Train Loss: 0.1319, Train Accuracy: 96.67%, Validation Loss: 0.0943, Validation Accuracy: 98.97%
Epoch [20/35] - Train Loss: 0.0642, Train Accuracy: 99.33%, Validation Loss: 0.0791, Validation Accuracy: 97.94%
Epoch [21/35] - Train Loss: 0.1029, Train Accuracy: 98.22%, Validation Loss: 0.0790, Validation Accuracy: 100.00%
Epoch [22/35] - Train Loss: 0.0827, Train Accuracy: 98.44%, Validation Loss: 0.0554, Validation Accuracy: 98.97%
Epoch [23/35] - Train Loss: 0.0755, Train Accuracy: 98.22%, Validation Loss: 0.1134, Validation Accuracy: 96.91%
Epoch [24/35] - Train Loss: 0.0432, Train Accuracy: 99.56%, Validation Loss: 0.0639, Validation Accuracy: 96.91%
Epoch [25/35] - Train Loss: 0.0541, Train Accuracy: 99.56%, Validation Loss: 0.0526, Validation Accuracy: 98.97%
Epoch [26/35] - Train Loss: 0.0428, Train Accuracy: 99.33%, Validation Loss: 0.0978, Validation Accuracy: 95.88%
Epoch [27/35] - Train Loss: 0.0911, Train Accuracy: 97.11%, Validation Loss: 0.1585, Validation Accuracy: 92.78%
Epoch [28/35] - Train Loss: 0.0628, Train Accuracy: 99.11%, Validation Loss: 0.0611, Validation Accuracy: 97.94%
Epoch [29/35] - Train Loss: 0.0678, Train Accuracy: 98.44%, Validation Loss: 0.0961, Validation Accuracy: 98.97%
Epoch [30/35] - Train Loss: 0.0537, Train Accuracy: 98.89%, Validation Loss: 0.0790, Validation Accuracy: 96.91%
Early stopping! Best validation loss: 0.0526

```

Test Loss: 0.1404, Test Accuracy: 94.85%



Here's a description of the results:

During the training phase of the hand gesture recognition model, 35 epochs were intended to be executed. However, due to the early stopping mechanism in place, the training was stopped at the 30th epoch to prevent overfitting, with the best validation loss recorded at 0.0526.

The training showed an improvement in accuracy over the epochs:

1. Training Performance: The training started with an accuracy of 34.44% in the first epoch, after then climbing to reach a high of 99.56% by the 25th epoch.
2. Validation Performance: For the validation set, accuracy began at 44.33% in the first epoch and increased significantly, touching a peak of 100% by the 21st epoch.
3. Test Performance: Once the training concluded, the model was evaluated on a test set, achieving an accuracy of 94.85% with a test loss of 0.1404.

Below, you'll find a visual comparison of loss and accuracy metrics between the training and validation sets and confusion matrix for gesture classes.



8 Challenges and limitations

One of the primary challenges encountered during the project was the inconsistency in the frame length across different gesture video segments in the dataset. This inconsistency posed a dimensionality challenge when interfacing with the neural network architecture, which expects a consistent input size. To address this, two potential solutions were identified:

Padding Technique: The first approach was to identify the video segment with the maximum frame length and then pad every other video to match this length. However, this method has a significant drawback. Padding essentially involves adding empty frames, which might not carry meaningful gesture information. This could introduce noise into the dataset and potentially mislead the model during training.

Downsampling: The alternative was to compute the average frame length across the entire dataset. For our specific dataset, this average was found to be 60 frames per second. Hence, all videos with frame lengths equal to or exceeding this average were selected. These selected videos were then downsampled to a consistent 60 frames. This method ensured that the videos retained meaningful gesture information without introducing artificial data. However, this approach resulted in the exclusion of some data samples, as gesture videos with frame lengths shorter than 60 were not considered. Moreover, the process of downsampling videos with a significantly higher number of frames to a consistent 60 frames can cause certain gesture videos to appear accelerated. This rapid pacing can introduce substantial variability in gesture dynamics, complicating the model's training due to the increased variations in gesture speed and motion. The choice between these strategies required careful consideration of their implications on the model's training and performance.

Another limitation arises from the need to maintain both spatial and temporal information within the dataset. This constraint allowed us to employ only rotation and translation as augmentation techniques, restricting our ability to diversify the dataset further.

9 Future work

Model Exploration:

While the current research utilized the ResNet3D_18 architecture, there's room for broadening the scope. Future endeavors should include training on a variety of neural network architectures to ascertain optimal performance and possibly uncover nuances unique to each model.

Emphasis on Data:

The quality and volume of the dataset remain paramount in deep learning tasks. Regardless of the sophistication of the chosen architecture, the underlying data can make or break model performance. Going forward, the objective should be to enhance the dataset, both in terms of diversity and quantity, to ensure robust and generalizable outcomes.

Preprocessing-Background Context:

The current approach extracts hand landmarks against a black background, removing the original context. A promising avenue for future research would be to train the model on hand landmarks superimposed on their original backgrounds. This could provide richer information and potentially improve the model's ability to generalize across diverse real-world scenarios.

Fine Tuning:

Owing to constraints in computational resources, the model's fine-tuning was not done greatly during this project. In future work, a more comprehensive fine-tuning process will be done to potentially enhance the model's performance.

References

1. Oudah M, Al-Naji A, Chahl J. Hand Gesture Recognition Based on Computer Vision: A Review of Techniques. Journal of Imaging
<https://doi.org/10.3390/jimaging6080073>
2. Hussain Z, Gimenez F, Yi D, Rubin D. Differential Data Augmentation Techniques for Medical Imaging Classification Tasks.
<https://www.ncbi.nlm.nih.gov/pmc/articles/PMC5977656/>
3. Bhatt, Dulari, Chirag Patel, Hardik Talsania, Jigar Patel, Rasmika Vaghela, Sharnil Pandya, Kirit Modi, and Hemant Ghayvat. 2021. "CNN Variants for Computer Vision: History, Architecture, Application, Challenges and Future Scope"
<https://doi.org/10.3390/electronics10202470>
4. Sharma, Sakshi; Singh, Sukhwinder. "ISL recognition system using integrated mobile-net and transfer learning method." Expert Systems with Applications
<https://doi.org/10.1016/j.eswa.2023.119772>
5. Kumar, Vijay; Alnuaim, Abeer; Zakariah (2022). "Human-Computer Interaction with Hand Gesture Recognition Using ResNet and MobileNet." Computational Intelligence and Neuroscience
<https://doi.org/10.1155/2022/8777355>
6. M. L. Amit, A. C. Fajardo and R. P. Medina, "Recognition of Real-Time Hand Gestures using Mediapipe Holistic Model and LSTM with MLP Architecture,"
<https://ieeexplore.ieee.org/abstract/document/10001800>
7. Abu Saleh Musa Miah, Md. Al Mehedi Hasan, And Jungpil Shin, "Dynamic Hand Gesture Recognition Using Multi-Branch Attention Based Graph and General Deep Learning Model," School of Computer Science and Engineering
<https://ieeexplore.ieee.org/stamp/stamp.jsp?arnumber=10012305>
8. Zhaofan Qiu, Ting Yao, Tao Mei; Proceedings of the IEEE International Conference on Computer Vision (ICCV)
https://www.microsoft.com/en-us/research/wp-content/uploads/2017/10/iccv_p3d_camera.pdf
9. Wang, S., Wang, K., Yang, T. *et al.* Improved 3D-ResNet sign language recognition algorithm with enhanced hand features.
<https://www.nature.com/articles/s41598-022-21636-z#citeas>
10. Bao W, Ma Z, Liang D, Yang X, Niu T. Pose ResNet: 3D Human Pose Estimation Based on Self-Supervision.
<https://www.mdpi.com/1424-8220/23/6/3057>

11. Hakan Bilen, Basura Fernando, Efstratios Gavves, Andrea Vedaldi; Dynamic Image neural network for action recognition Networks for Action Recognition

https://www.researchgate.net/publication/319306777_Action_Recognition_with_Dynamic_Image_Networks

12. T. Alshalali and D. Josyula, "Fine-Tuning of Pre-Trained Deep Learning Models with Extreme Learning Machine," 2018 International Conference on Computational Science and Computational Intelligence (CSCI)

<https://ieeexplore.ieee.org/abstract/document/8947855>